

PySpark Basics

Learn Python for data science **Interactively** at www.DataCamp.com



Spark

PySpark is the Spark Python API that exposes the Spark programming model to Python



Initializing Spark

SparkContext

```
>>> from pyspark import SparkContext
>>> sc = SparkContext(master = 'local[2]')
```

Inspect SparkContext

>>> sc.version	Retrieve SparkContext version
>>> sc.pythonVer	Retrieve Python version
>>> sc.master	Master URL to connect to
>>> str(sc.sparkHome)	Path where Spark is installed on worker nodes
>>> str(sc.sparkUser())	Retrieve name of the Spark User running SparkContext
>>> sc.appName	Return application name
>>> sc.applicationId	Retrieve application ID
>>> sc.defaultParallelism	Return default level of parallelism
>>> sc.defaultMinPartitions	Default minimum number of partitions for RDDs

Configuration

```
>>> from pyspark import SparkConf, SparkContext
>>> conf = (SparkConf()
            .setMaster("local")
            .setAppName("My app")
            .set("spark.executor.memory", "1g"))
>>> sc = SparkContext(conf = conf)
```

Using The Shell

In the PySpark shell, a special interpreter-aware SparkContext is already created in the variable called `sc`.

```
$ ./bin/spark-shell --master local[2]
$ ./bin/pyspark --master local[4] --py-files code.py
```

Set which master the context connects to with the `--master` argument, and add Python `.zip`, `.egg` or `.py` files to the runtime path by passing a comma-separated list to `--py-files`.

Loading Data

Parallelized Collections

```
>>> rdd = sc.parallelize([('a',7), ('a',2), ('b',2)])
>>> rdd2 = sc.parallelize([('a',2), ('d',1), ('b',1)])
>>> rdd3 = sc.parallelize(range(100))
>>> rdd4 = sc.parallelize([("a",["x","y","z"]),
                          ("b",["p","r"])]])
```

External Data

Read either one text file from HDFS, a local file system or or any Hadoop-supported file system URI with `textFile()`, or read in a directory of text files with `wholeTextFiles()`.

Basic Information

>>> rdd.getNumPartitions()	List the number of partitions
>>> rdd.count()	Count RDD instances
3	
>>> rdd.countByKey()	Count RDD instances by key
defaultdict(<type 'int'>, {'a':2, 'b':1})	
>>> rdd.countByValue()	Count RDD instances by value
defaultdict(<type 'int'>, {'b':2}:1, ('a',2):1, ('a',7):1)	
>>> rdd.collectAsMap()	Return (key,value) pairs as a dictionary
{'a': 2, 'b': 2}	Sum of RDD elements
>>> rdd3.sum()	
4950	
>>> sc.parallelize([]).isEmpty()	Check whether RDD is empty
True	

Summary

>>> rdd3.max()	Maximum value of RDD elements
99	
>>> rdd3.min()	Minimum value of RDD elements
0	
>>> rdd3.mean()	Mean value of RDD elements
49.5	
>>> rdd3.stdev()	Standard deviation of RDD elements
28.866070047722118	
>>> rdd3.variance()	Compute variance of RDD elements
833.25	
>>> rdd3.histogram(3)	Compute histogram by bins
[[0, 33, 66, 99], [33, 33, 34]]	
>>> rdd3.stats()	Summary statistics (count, mean, stdev, max & min)

Applying Functions

>>> rdd.map(lambda x: x+(x[1],x[0]))	Apply a function to each RDD element
.collect()	
[('a',7,7,'a'), ('a',2,2,'a'), ('b',2,2,'b')]	
>>> rdd5 = rdd.flatMap(lambda x: x+(x[1],x[0]))	Apply a function to each RDD element and flatten the result
>>> rdd5.collect()	
[('a',7,7,'a','a',2,2,'a','b',2,2,'b')]	
>>> rdd4.flatMapValues(lambda x: x)	Apply a flatMap function to each (key,value) pair of rdd4 without changing the keys
.collect()	
[('a','x'), ('a','y'), ('a','z'), ('b','p'), ('b','r')]	

Selecting Data

Getting	
>>> rdd.collect()	Return a list with all RDD elements
[('a', 7), ('a', 2), ('b', 2)]	
>>> rdd.take(2)	Take first 2 RDD elements
[('a', 7), ('a', 2)]	
>>> rdd.first()	Take first RDD element
('a', 7)	
>>> rdd.top(2)	Take top 2 RDD elements
[('b', 2), ('a', 7)]	
Sampling	
>>> rdd3.sample(False, 0.15, 81).collect()	Return sampled subset of rdd3
[3, 4, 27, 31, 40, 41, 42, 43, 60, 76, 79, 80, 86, 97]	
Filtering	
>>> rdd.filter(lambda x: "a" in x)	Filter the RDD
.collect()	
[('a',7), ('a',2)]	
>>> rdd5.distinct().collect()	Return distinct RDD values
[('a',2,'b',7)]	
>>> rdd.keys().collect()	Return (key,value) RDD's keys
[('a', 'a', 'b')]	

Iterating

```
>>> def g(x): print(x)
>>> rdd.foreach(g)
```

Apply a function to all RDD elements

Reducing

```
>>> rdd.reduceByKey(lambda x,y : x+y)
.collect()
[('a',9), ('b',2)]
>>> rdd.reduce(lambda a, b: a + b)
('a',7, 'a',2, 'b',2)

Grouping by
>>> rdd3.groupBy(lambda x: x % 2)
.mapValues(list)
.collect()
>>> rdd.groupByKey()
.mapValues(list)
.collect()
[('a', [7,2]), ('b', [2])]
```

Aggregating

```
>>> seqOp = (lambda x,y: (x[0]+y,x[1]+1))
>>> combOp = (lambda x,y: (x[0]+y[0],x[1]+y[1]))
>>> rdd3.aggregate((0,0),seqOp,combOp)
(4950,100)
>>> rdd.aggregateByKey((0,0),seqOp,combOp)
.collect()
[('a', (9,2)), ('b', (2,1))]
>>> rdd3.fold(0,add)
4950
>>> rdd.foldByKey(0, add)
.collect()
[('a',9), ('b',2)]
>>> rdd3.keyBy(lambda x: x+x)
.collect()
```

Merge the rdd values for each key

Merge the rdd values

Return RDD of grouped values

Group rdd by key

Aggregate RDD elements of each partition and then the results

Aggregate values of each RDD key

Aggregate the elements of each partition, and then the results

Merge the values for each key

Create tuples of RDD elements by applying a function

Mathematical Operations

>>> rdd.subtract(rdd2)	Return each rdd value not contained in rdd2
.collect()	
[('b',2), ('a',7)]	
>>> rdd2.subtractByKey(rdd)	Return each (key,value) pair of rdd2 with no matching key in rdd
.collect()	
[('d', 1)]	
>>> rdd.cartesian(rdd2).collect()	Return the Cartesian product of rdd and rdd2

Sort

>>> rdd2.sortBy(lambda x: x[1])	Sort RDD by given function
.collect()	
[('d',1), ('b',1), ('a',2)]	
>>> rdd2.sortByKey()	Sort (key, value) RDD by key
.collect()	
[('a',2), ('b',1), ('d',1)]	

Repartitioning

```
>>> rdd.repartition(4)
>>> rdd.coalesce(1)
```

New RDD with 4 partitions

Decrease the number of partitions in the RDD to 1

Saving

```
>>> rdd.saveAsTextFile("rdd.txt")
>>> rdd.saveAsHadoopFile("hdfs://namenodehost/parent/child",
                        'org.apache.hadoop.mapred.TextOutputFormat')
```

Stopping SparkContext

```
>>> sc.stop()
```

Execution

```
$ ./bin/spark-submit examples/src/main/python/pi.py
```